

day1_eda

```
churn <- read_csv(here('data/raw', 'Churn_Modelling.csv'))  
head(churn)
```

```
# A tibble: 6 × 14  
RowNumber CustomerId Surname CreditScore Geography Gender  
Age Tenure Balance  
<dbl>      <dbl> <chr>          <dbl> <chr>    <chr>  
<dbl> <dbl>   <dbl>  
1       1   15634602 Hargra...        619 France Female  
42      2     0  
2       2   15647311 Hill           608 Spain  Female  
41      1  83808.  
3       3   15619304 Onio          502 France Female  
42      8  159661.  
4       4   15701354 Boni          699 France Female  
39      1     0  
5       5   15737888 Mitchel...      850 Spain  Female  
43      2 125511.  
6       6   15574012 Chu            645 Spain  Male  
44      8 113756.  
# i 5 more variables: NumOfProducts <dbl>, HasCrCard <dbl>,  
#   IsActiveMember <dbl>, EstimatedSalary <dbl>, Exited <dbl>
```

```
nrow(churn)
```

```
[1] 10000
```

```
ncol(churn)
```

```
[1] 14
```

```
glimpse(churn)
```

```
Rows: 10,000  
Columns: 14  
$ RowNumber      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,  
13, 14, 15, 16,...  
$ CustomerId     <dbl> 15634602, 15647311, 15619304,  
15701354, 15737888, 1557...  
$ Surname        <chr> "Hargrave", "Hill", "Onio", "Boni",  
"Mitchell", "Chu", ...  
$ CreditScore    <dbl> 619, 608, 502, 699, 850, 645, 822,
```

```

376, 501, 684, 528, ...
$ Geography      <chr> "France", "Spain", "France", "France",
"Spain", "Spain...
$ Gender         <chr> "Female", "Female", "Female",
"Female", "Female", "Mal...
$ Age            <dbl> 42, 41, 42, 39, 43, 44, 50, 29, 44,
27, 31, 24, 34, 25...
$ Tenure          <dbl> 2, 1, 8, 1, 2, 8, 7, 4, 4, 2, 6, 3,
10, 5, 7, 3, 1, 9, ...
$ Balance         <dbl> 0.00, 83807.86, 159660.80, 0.00,
125510.82, 113755.78, ...
$ NumOfProducts   <dbl> 1, 1, 3, 2, 1, 2, 2, 4, 2, 1, 2, 2, 2,
2, 2, 2, 1, 2, ...
$ HasCrCard       <dbl> 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1,
0, 1, 0, 1, ...
$ IsActiveMember  <dbl> 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0,
0, 1, 1, 0, 1, ...
$ EstimatedSalary <dbl> 101348.88, 112542.58, 113931.57,
93826.63, 79084.10, 1...
$ Exited          <dbl> 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 0, 1, 0, ...

```

`colSums(is.na(churn))`

	RowNumber	CustomerId	Surname
CreditScore	Geography		
	0	0	0
0	0		
	Gender	Age	Tenure
Balance	NumOfProducts		
	0	0	0
0	0		
	HasCrCard	IsActiveMember	EstimatedSalary
Exited			
	0	0	0
0			

`summary(churn)`

	RowNumber	CustomerId	Surname
CreditScore			
Min. : 1	Min. :15565701	Length:10000	Min.
:350.0			
1st Qu.: 2501	1st Qu.:15628528	Class :character	1st
Qu.:584.0			
Median : 5000	Median :15690738	Mode :character	Median
:652.0			

Mean : 5000	Mean : 15690941	Mean
: 650.5		
3rd Qu.: 7500	3rd Qu.: 15753234	3rd
Qu.: 718.0		
Max. : 10000	Max. : 15815690	Max.
: 850.0		
Geography	Gender	Age
Tenure		
Length: 10000	Length: 10000	Min. : 18.00
: 0.000		Min.
Class : character	Class : character	1st Qu.: 32.00
Qu.: 3.000		1st
Mode : character	Mode : character	Median : 37.00
: 5.000		Median
		Mean : 38.92
		Mean
: 5.013		
		3rd Qu.: 44.00
		3rd
Qu.: 7.000		
		Max. : 92.00
		Max.
: 10.000		
Balance	NumOfProducts	HasCrCard
IsActiveMember		
Min. : 0	Min. : 1.00	Min. : 0.0000
: 0.0000		Min.
1st Qu.: 0	1st Qu.: 1.00	1st Qu.: 0.0000
Qu.: 0.0000		1st
Median : 97199	Median : 1.00	Median : 1.0000
: 1.0000		Median
Mean : 76486	Mean : 1.53	Mean : 0.7055
: 0.5151		Mean
3rd Qu.: 127644	3rd Qu.: 2.00	3rd Qu.: 1.0000
Qu.: 1.0000		3rd
Max. : 250898	Max. : 4.00	Max. : 1.0000
: 1.0000		Max.
EstimatedSalary	Exited	
Min. : 11.58	Min. : 0.0000	
1st Qu.: 51002.11	1st Qu.: 0.0000	
Median : 100193.91	Median : 0.0000	
Mean : 100090.24	Mean : 0.2037	
3rd Qu.: 149388.25	3rd Qu.: 0.0000	
Max. : 199992.48	Max. : 1.0000	

Churn Rate: 20.4%

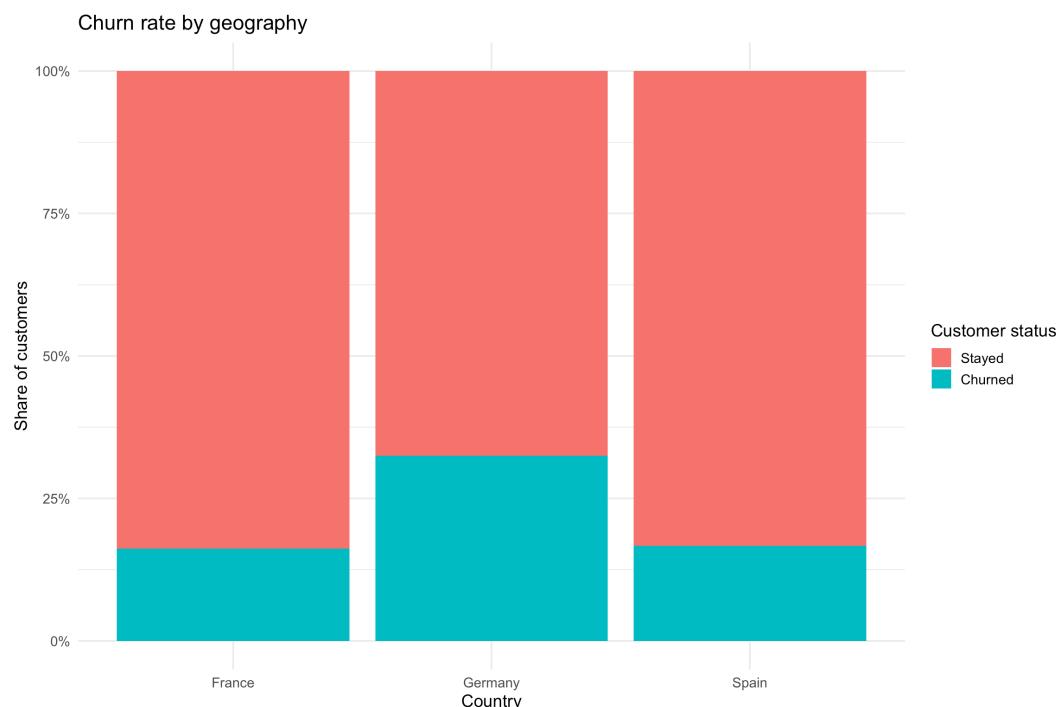
```
churn_rate <- churn |>
  count(Exited) |>
  mutate(pct = n/sum(n)*100)
```

churn_rate

```
# A tibble: 2 × 3
  Exited      n    pct
  <dbl> <int> <dbl>
1     0    7963  79.6
2     1   2037  20.4
```

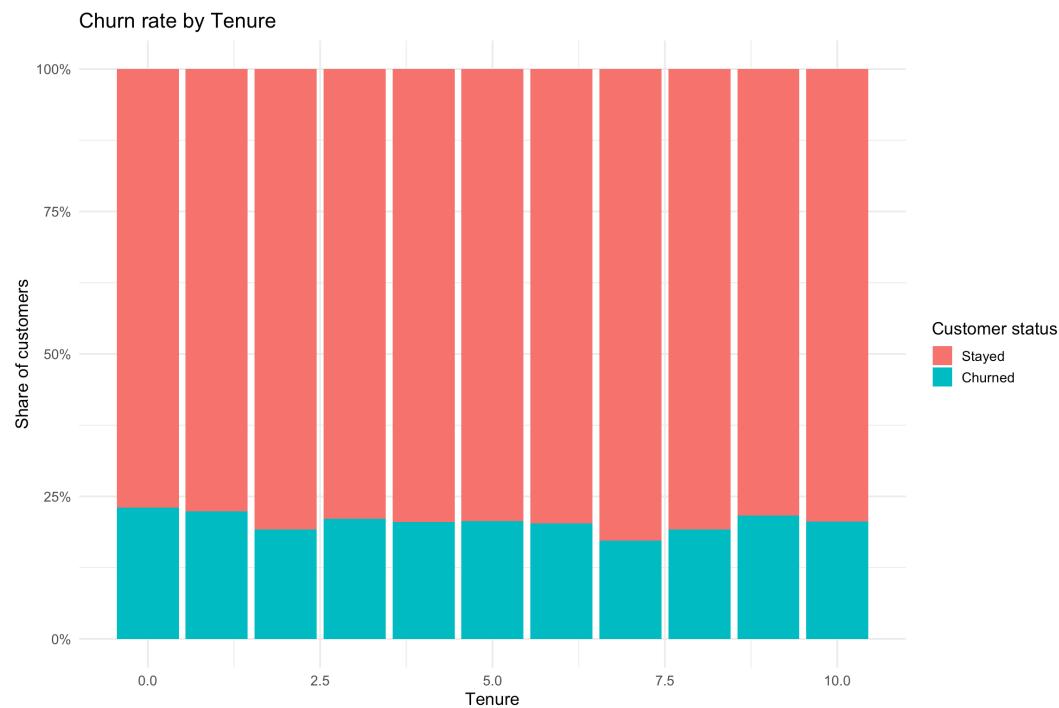
~80% of the people are happy and have stayed for this period.

```
ggplot(churn, aes(x = Geography, fill = factor(Exited))) +
  geom_bar(position = "fill") +
  scale_fill_discrete(
    name = "Customer status",
    labels = c("0" = "Stayed", "1" = "Churned")
  ) +
  scale_y_continuous(labels = scales::label_percent(accuracy =
  labs(
    title = "Churn rate by geography",
    x = "Country",
    y = "Share of customers"
  ) +
  theme_minimal(base_size = 14)
```



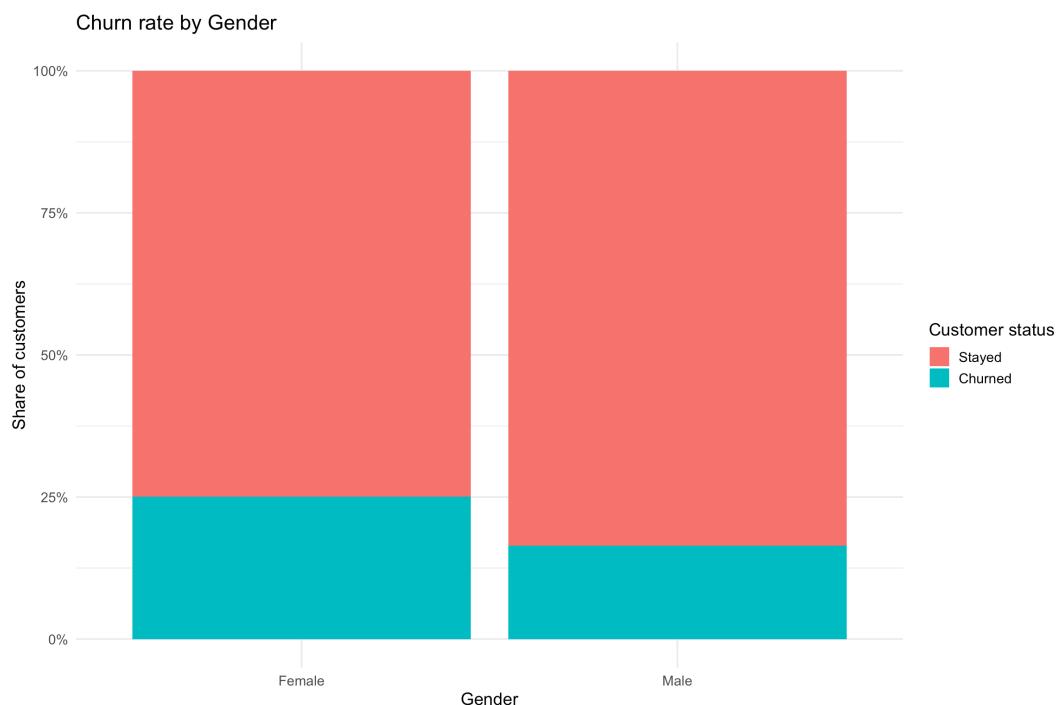
Germany has the maximum churned share of customers close to 30% , both France and Spain have closely equal share of customer who have churned (~15%).

```
ggplot(churn, aes(x = Tenure, fill = factor(Exited))) +  
  geom_bar(position = "fill") +  
  scale_fill_discrete(  
    name = "Customer status",  
    labels = c("0" = "Stayed", "1" = "Churned"))  
  ) +  
  scale_y_continuous(labels = scales::label_percent(accuracy =  
    labs(  
      title = "Churn rate by Tenure",  
      x = "Tenure",  
      y = "Share of customers"  
    ) +  
    theme_minimal(base_size = 14)
```



we see a uniform share of customers through all the tenure period, which mean customers remain loyal irrespective of whether being new or old(~18 to 22%).

```
ggplot(churn, aes(x = Gender, fill = factor(Exited))) +  
  geom_bar(position = "fill") +  
  scale_fill_discrete(  
    name = "Customer status",  
    labels = c("0" = "Stayed", "1" = "Churned"))  
  ) +  
  scale_y_continuous(labels = scales::label_percent(accuracy =  
    labs(  
      title = "Churn rate by Gender",  
      x = "Gender",  
      y = "Share of customers"  
    ) +  
    theme_minimal(base_size = 14)
```



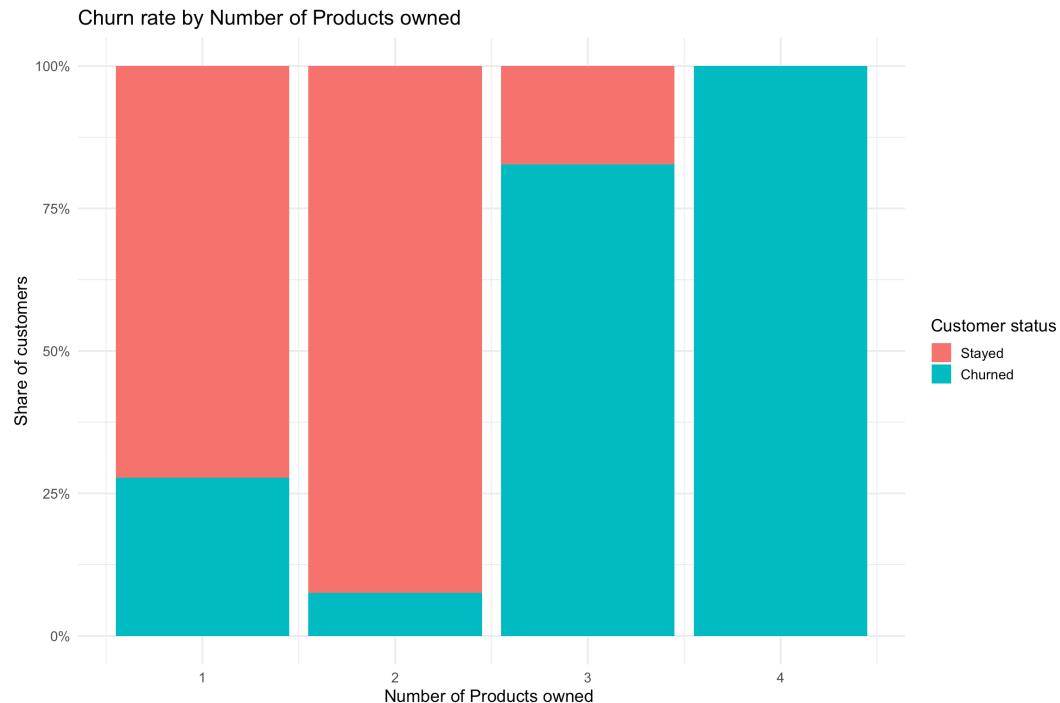
Female customer has churned more than male, with 25% for female and ~ 18% for male customers.

```
ggplot(churn, aes(x = NumOfProducts, fill = factor(Exited))) +  
  geom_bar(position = "fill") +  
  scale_fill_discrete(  
    name = "Customer status",  
    labels = c("0" = "Stayed", "1" = "Churned"))  
  ) +  
  scale_y_continuous(labels = scales::label_percent(accuracy =
```

```

labs(
  title = "Churn rate by Number of Products owned",
  x = "Number of Products owned",
  y = "Share of customers"
) +
theme_minimal(base_size = 14)

```

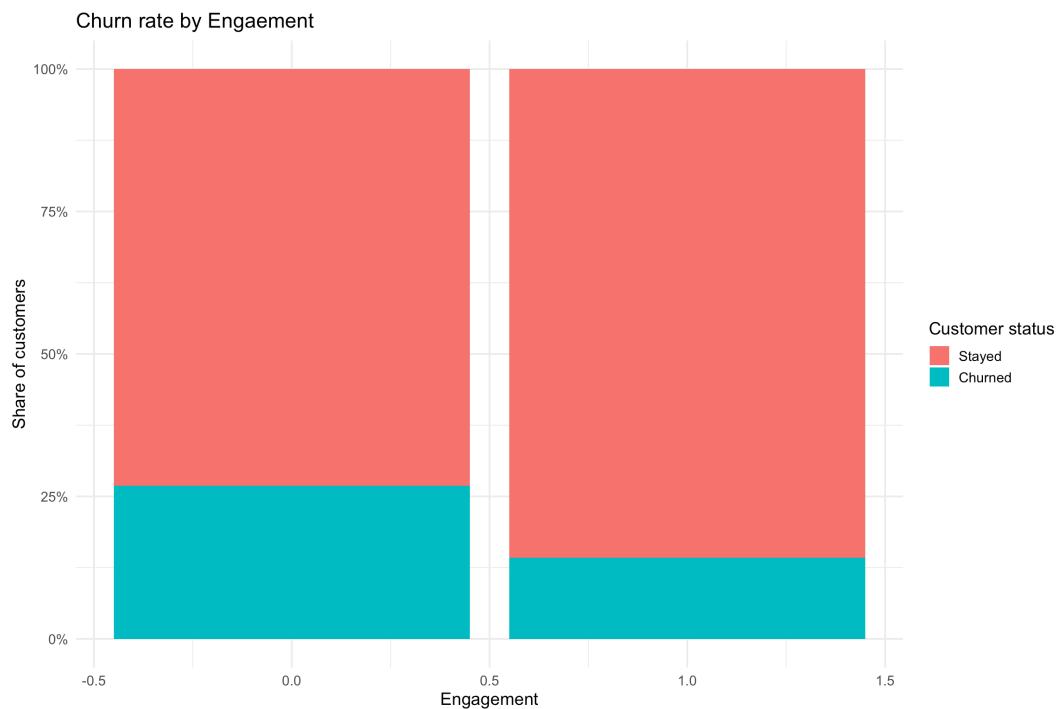


People with 3 and 4 products have churned most, with 4 product being 100 %. Its a huge driver for customer exiting.

```

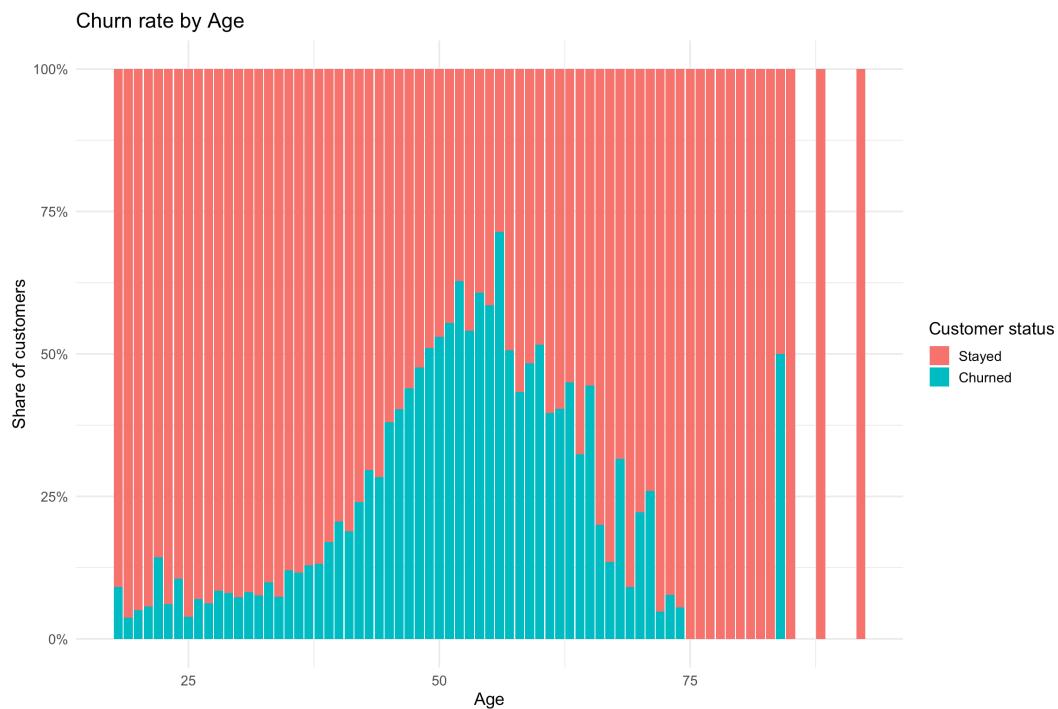
ggplot(churn, aes(x = IsActiveMember, fill = factor(Exited))) +
  geom_bar(position = "fill") +
  scale_fill_discrete(
    name = "Customer status",
    labels = c("0" = "Stayed", "1" = "Churned")
  ) +
  scale_y_continuous(labels = scales::label_percent(accuracy =
  labs(
    title = "Churn rate by Engagement",
    x = "Engagement",
    y = "Share of customers"
  ) +
  theme_minimal(base_size = 14)

```



People who are not active have churned more than people who are active.

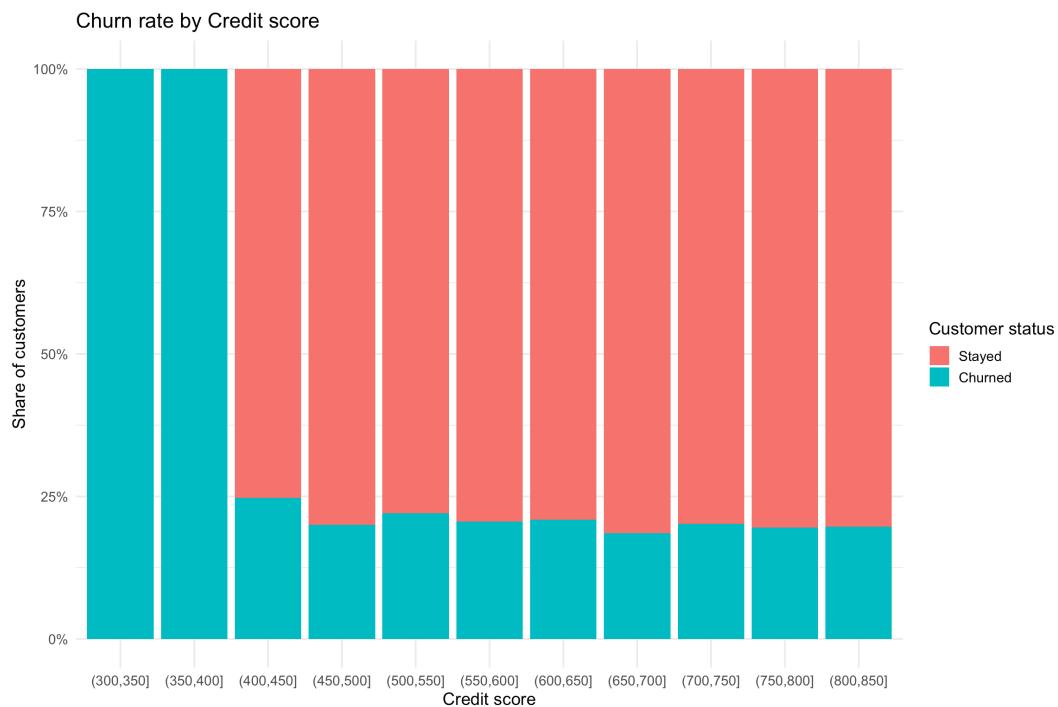
```
ggplot(churn, aes(x = Age, fill = factor(Exited))) +  
  geom_bar(position = "fill") +  
  scale_fill_discrete(  
    name = "Customer status",  
    labels = c("0" = "Stayed", "1" = "Churned"))  
  +  
  scale_y_continuous(labels = scales::label_percent(accuracy =  
    labs(  
      title = "Churn rate by Age",  
      x = "Age",  
      y = "Share of customers"  
    ) +  
    theme_minimal(base_size = 14)
```



it seems to normal distributed, with mean close 50-55 yrs, which means people close to 50 to 55 yrs has churned most. it seems to have a outlier at ~80yrs, maybe due to people's death.

```
churn_df <- churn|>
  mutate(credit_bin = cut(CreditScore, breaks=seq(300,850,50)))

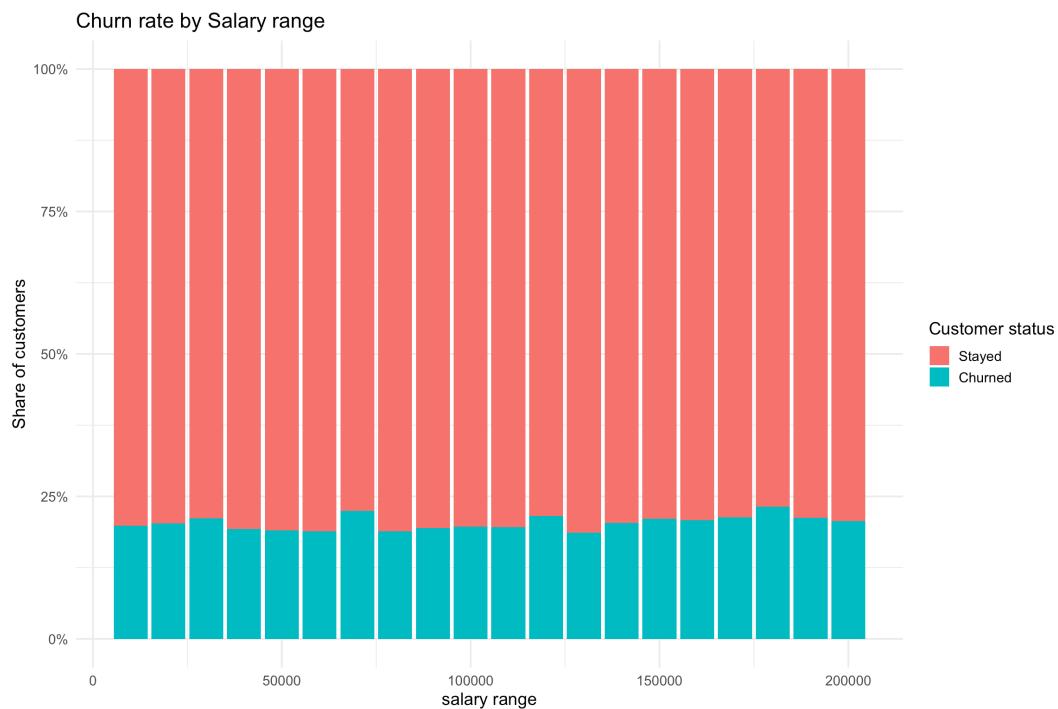
ggplot(churn_df, aes(x = credit_bin, fill = factor(Exited))) +
  geom_bar(position = "fill") +
  scale_fill_discrete(
    name   = "Customer status",
    labels = c("0" = "Stayed", "1" = "Churned")
  ) +
  scale_y_continuous(labels = scales::label_percent(accuracy =
  labs(
    title = "Churn rate by Credit score",
    x = "Credit score",
    y = "Share of customers"
  ) +
  theme_minimal(base_size = 14)
```



When the customers credit score is less than 400, their churning rate is very high close to 100% percent. the chruning rate decrease gradually as the credit score increases.

```
churn_df <- churn|>
  mutate(salary_bin = ceiling(EstimatedSalary/10000)*10000)

ggplot(churn_df, aes(x = salary_bin, fill = factor(Exited))) +
  geom_bar(position = "fill") +
  scale_fill_discrete(
    name   = "Customer status",
    labels = c("0" = "Stayed", "1" = "Churned")
  ) +
  scale_y_continuous(labels = scales::label_percent(accuracy =
  labs(
    title = "Churn rate by Salary range",
    x = "salary range",
    y = "Share of customers"
  ) +
  theme_minimal(base_size = 14)
```



we see a uniform share of customers through all the Salary range, which mean customers remain loyal irrespective of their salary (~20 to 25%).

```
library(corrplot)
```

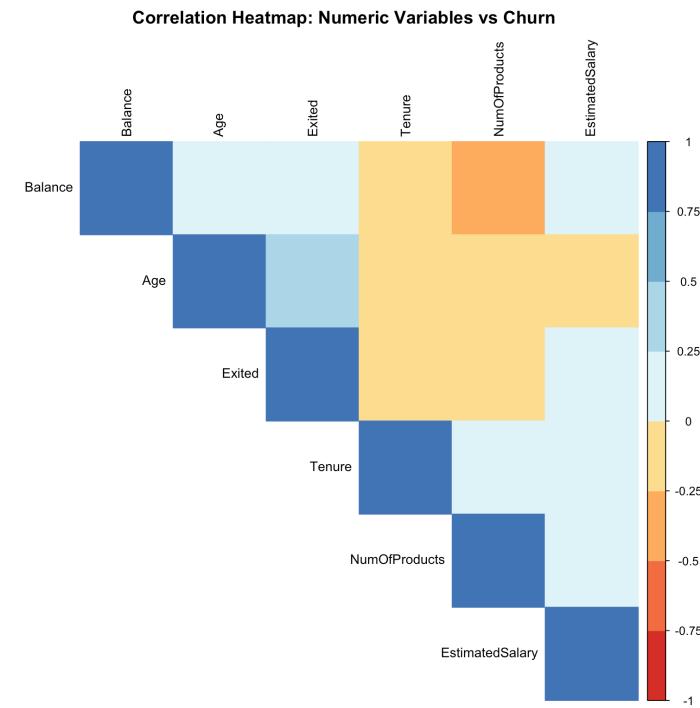
corrplot 0.95 loaded

```
library(RColorBrewer)

# Select numeric vars for correlation
num_vars <- churn %>%
  select(Age, Tenure, Balance, NumOfProducts, EstimatedSalary,
  cor(use = "complete.obs")

# Heatmap
corrplot(num_vars,
  method = "color",
  type = "upper",
  order = "hclust",
  tl.cex = 0.9,
  tl.col = "black",
  col = brewer.pal(n = 8, name = "RdYlBu"),
```

```
title = "Correlation Heatmap: Numeric Variables vs Churn"
mar = c(0,0,2,0))
```



"hypotheses for modeling" summary:

- Higher churn in Germany.
- Higher churn in females and mid-age customers.
- Very high churn when credit score < 400.
- Very high churn when people have more than 3 products.
- Higher churn in people who are not active.
- Tenure and salary look relatively flat vs churn.
- no serious multicollinearity problem among these numeric features, and only a weak-to-moderate relationship between the predictors and churn